



▼ Задание 3

В этом задании вам предстоит попробовать несколько методов, используемых в задаче классификации, а также понять насколько хорошо модель понимает смысл слов и какие слова в примере влияют на результат.

```
Collecting torchtext==0.8.1
  Downloading https://files.pythonhosted.org/packages/13/80/046f0691b296e755ae884df3c/ | ██████████ 7.0MB 3.4MB/s
Collecting torch==1.7.1
  Downloading https://files.pythonhosted.org/packages/90/5d/095ddddd91c8a769a68c791cf/ | ██████████ 776.8MB 23kB/s
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from torch)
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from torch)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torch)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch)
Requirement already satisfied: urllib3!=1.25.0,!<1.25.1,<1.26,>=1.21.1 in /usr/local,
```

```
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages
ERROR: torchvision 0.9.0+cu101 has requirement torch==1.8.0, but you'll have torch 1
Installing collected packages: torch, torchtext
  Found existing installation: torch 1.8.0+cu101
    Uninstalling torch-1.8.0+cu101:
      Successfully uninstalled torch-1.8.0+cu101
  Found existing installation: torchtext 0.9.0
    Uninstalling torchtext-0.9.0:
      Successfully uninstalled torchtext-0.9.0
Successfully installed torch-1.7.1 torchtext-0.8.1
```

```
!pip install torch==1.7.1+cu101 torchvision==0.8.2+cu101 torchaudio==0.7.2 -f
```

Usage:

```
pip3 install [options] <requirement specifier> [package-index-options] ...
pip3 install [options] -r <requirements file> [package-index-options] ...
pip3 install [options] [-e] <vcs project url> ...
pip3 install [options] [-e] <local project path> ...
pip3 install [options] <archive url/path> ...
```

-f option requires 1 argument

```
import pandas as pd
import numpy as np
import torch
```

```
from torchtext import datasets
```

```
from torchtext.data import Field, LabelField
from torchtext.data import BucketIterator
```

```
from torchtext.vocab import Vectors, GloVe
```

```
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import random
from tqdm.autonotebook import tqdm
```

В этом задании мы будем использовать библиотеку torchtext. Она довольно проста в использовании и поможет нам сконцентрироваться на задаче, а не на написании DataLoader-a.

```
TEXT = Field(sequential=True, lower=True, include_lengths=True) # Поле текста
LABEL = LabelField(dtype=torch.float) # Поле метки
```

```
/usr/local/lib/python3.7/dist-packages/torchtext/data/field.py:150: UserWarning: Field
warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please
/usr/local/lib/python3.7/dist-packages/torchtext/data/field.py:150: UserWarning: Label
warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please
```



```

113, 113, 113, 113, 113, 113, 113, 113], device='cuda:0'))
tensor([1., 1., 1., 1., 1., 1., 0., 0., 0., 0., 1., 0., 0., 1., 0., 1., 0., 1.,
        0., 1., 1., 0., 0., 1., 0., 1., 0., 1., 0., 1., 1., 1., 0., 1., 0., 1.,
        0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 0., 0.,
        0., 1., 1., 1., 1., 0., 1., 1., 1., 0.], device='cuda:0')
114

```

▼ RNN

Для начала попробуем использовать рекуррентные нейронные сети. На семинаре вы познакомились с GRU, вы можете также попробовать LSTM. Можно использовать для классификации как `hidden_state`, так и `output` последнего токена.

```

class RNNBaseline(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim, n_layers,
                  bidirectional, dropout, pad_idx):

        super().__init__()

        self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx = pad_idx)

        self.rnn = nn.LSTM(input_size=embedding_dim, hidden_size=hidden_dim, num_layers=n_

        self.fc = nn.Linear(n_layers*hidden_dim, output_dim)  # YOUR CODE GOES HERE

        self.dropout = nn.Dropout(dropout)

    def forward(self, text, text_lengths):

        #text = [sent len, batch size]

        embedded = self.embedding(text)

        #embedded = [sent len, batch size, emb dim]

        #pack sequence
        packed_embedded = nn.utils.rnn.pack_padded_sequence(embedded, text_lengths)
        # print("packed_embedded=", packed_embedded.data.shape)

        # cell arg for LSTM, remove for GRU
        packed_output, (hidden, cell) = self.rnn(packed_embedded)

        #unpack sequence
        output, output_lengths = nn.utils.rnn.pad_packed_sequence(packed_output)

        #output = [sent len, batch size, hid dim * num directions]
        #output over padding tokens are zero tensors

        #hidden = [num layers * num directions, batch size, hid dim]
        #cell = [num layers * num directions, batch size, hid dim]

```

```
#concat the final forward (hidden[-2,:,:]) and backward (hidden[-1,:,:]) hidden la
#and apply dropout

hiddens = torch.cat((hidden[-2,:,:], hidden[-1,:,:]), 1) #None # YOUR CODE GOES H
hiddens = self.dropout(hiddens)

return self.fc(hiddens)
```

Поиграйтесь с гиперпараметрами

```
vocab_size = len(TEXT.vocab)
emb_dim = 100
hidden_dim = 256
output_dim = 1
n_layers = 2
bidirectional = True
dropout = 0.8
PAD_IDX = TEXT.vocab.stoi[TEXT.pad_token]
patience=5
```

```
model = RNNBaseline(
    vocab_size=vocab_size,
    embedding_dim=emb_dim,
    hidden_dim=hidden_dim,
    output_dim=output_dim,
    n_layers=n_layers,
    bidirectional=bidirectional,
    dropout=dropout,
    pad_idx=PAD_IDX
)
```

```
model = model.to(device)
```

```
opt = torch.optim.Adam(model.parameters())
loss_func = nn.BCEWithLogitsLoss()
loss_func1 = nn.MSELoss(reduction="sum")
```

```
max_epochs = 20
```

```
x = torch.tensor([[1, 2, 3, 4]])
torch.unsqueeze(x, 2)
```

```
tensor([[[1],
         [2],
         [3],
         [4]]])
```

Обучите сетку! Используйте любые вам удобные инструменты, Catalyst, PyTorch Lightning или свои велосипеды.

```

def list_pred(list_pred):
    a = []
    for i in range(len(list_pred)):
        if i > 0.5:
            a.append(1)
        else:
            a.append(0)

    return a

import numpy as np
from sklearn.metrics import accuracy_score

min_loss = np.inf

cur_patience = 0

for epoch in range(1, max_epochs + 1):

    all_preds_train = []
    all_label_train = []
    all_preds_val = []
    all_label_val = []
    correct_train = 0
    correct_val = 0
    num_obj_train = 0
    num_obj_val = 0

    train_loss = 0.0
    model.train()
    pbar = tqdm(enumerate(train_iter), total=len(train_iter), leave=False)
    pbar.set_description(f"Epoch {epoch}")
    for it, batch in pbar:
        opt.zero_grad()
        emb = batch.text[0].to(device)
        len_emb = batch.text[1].cpu()
        label = torch.unsqueeze(batch.label, 1).to(device)
        preds = model(emb, len_emb)
        loss = loss_func(preds, label)
        loss.backward()
        train_loss += loss
        opt.step()

        preds_r = np.where(preds.cpu() > 0, 1, 0)
        label_r = np.where(label.cpu() > 0, 1, 0)
        c = preds_r == label_r
        correct_train += np.count_nonzero(c)
        num_obj_train += len(batch.label)

    trn = np.sum(preds_r, axis=1)
    lbl = np.sum(label_r, axis=1)
    trn = trn.tolist()
    lbl = lbl.tolist()
    all_preds_train.append(trn)

```

```
all_preds_train.append(preds)
all_label_train.append(lbl)
```

```
#YOUR CODE GOES HERE
```

```
train_loss /= len(train_iter)
val_loss = 0.0
model.eval()
with torch.no_grad():
    pbar = tqdm(enumerate(valid_iter), total=len(valid_iter), leave=False)
    pbar.set_description(f"Epoch {epoch}")
    for it, batch in pbar:
        emb = batch.text[0].to(device)
        len_emb = batch.text[1].cpu()
        label = torch.unsqueeze(batch.label, 1).to(device)
        preds = model(emb, len_emb)

        preds_r = np.where(preds.cpu() > 0, 1, 0)
        label_r = np.where(label.cpu() > 0, 1, 0)
        c = preds_r == label_r
        correct_val += np.count_nonzero(c)
        num_obj_val += len(batch.label)

    val_loss += loss_func(preds.data, label)

    val = np.sum(preds_r, axis=1)
    lbl_val = np.sum(label_r, axis=1)
    val = val.tolist()
    lbl_val = lbl_val.tolist()
    all_preds_val.append(val)
    all_label_val.append(lbl_val)
    # YOUR CODE GOES HERE
val_loss /= len(valid_iter)
if val_loss < min_loss:
    min_loss = val_loss
    best_model = model.state_dict()
else:
    cur_patience += 1
    if cur_patience == patience:
        cur_patience = 0
        break
```

```
print('Epoch: {}, Training Loss: {}, Validation Loss: {}, Training accuracy: {}, Valid
model.load_state_dict(best_model)
```

```
/usr/local/lib/python3.7/dist-packages/torchtext/data/batch.py:23: UserWarning: Batch
warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please
Epoch: 1, Training Loss: 0.6482466459274292, Validation Loss: 0.5980693101882935, Tra
Epoch: 2, Training Loss: 0.5436175465583801, Validation Loss: 0.5214125514030457, Tra
Epoch: 3, Training Loss: 0.5754403471946716, Validation Loss: 0.5855757594108582, Tra
Epoch: 4, Training Loss: 0.5563846230506897, Validation Loss: 0.6016130447387695, Tra
Epoch: 5, Training Loss: 0.3923629820346832, Validation Loss: 0.41213083267211914, Tr
Epoch: 6, Training Loss: 0.22027228772640228, Validation Loss: 0.4108792543411255, Tr
Epoch: 7, Training Loss: 0.11888790875673294, Validation Loss: 0.43305104970932007, 1
Epoch: 8, Training Loss: 0.06524424254894257, Validation Loss: 0.5048354268074036, Tr
<All keys matched successfully>
```

Посчитайте f1-score вашего классификатора на тестовом датасете.


Ответ:

```
flat_preds_list = [item for sublist in all_preds_val for item in sublist]
flat_label_list = [item for sublist in all_label_val for item in sublist]
```

```
from sklearn.metrics import f1_score
f1_score(flat_label_list, flat_preds_list, average='weighted')
```

0.8418184572607531

▼ CNN



Для классификации текстов также часто используют сверточные нейронные сети. Идея в том, что как правило sentiment содержат словосочетания из двух-трех слов, например "очень хороший фильм" или "невероятная скука". Проходясь сверткой по этим словам мы получим какой-то большой скор и выхватим его с помощью MaxPool. Далее идет обычная полносвязная сетка. Важный момент: свертки применяются не последовательно, а параллельно. Давайте попробуем!

```
TEXT = Field(sequential=True, lower=True, batch_first=True) # batch_first тк мы используем
LABEL = LabelField(batch_first=True, dtype=torch.float)
```

```
train, tst = datasets.IMDB.splits(TEXT, LABEL)
trn, vld = train.split(random_state=random.seed(SEED))
```



```
TEXT.build_vocab(trn)
LABEL.build_vocab(trn)
```

```
device = "cuda" if torch.cuda.is_available() else "cpu"
```

```
/usr/local/lib/python3.7/dist-packages/torchtext/data/field.py:150: UserWarning: Field
warnings.warn('{}' class will be retired soon and moved to torchtext.legacy. Please
/usr/local/lib/python3.7/dist-packages/torchtext/data/field.py:150: UserWarning: Label
warnings.warn('{}' class will be retired soon and moved to torchtext.legacy. Please
/usr/local/lib/python3.7/dist-packages/torchtext/data/example.py:78: UserWarning: Example
warnings.warn('Example class will be retired soon and moved to torchtext.legacy. Pl
```

```
train_iter, val_iter, test_iter = BucketIterator.splits(
    (trn, vld, tst),
    batch_sizes=(128, 256, 256),
    sort=False,
    sort_key= lambda x: len(x.src),
    sort_within_batch=False,
    device=device,
    repeat=False,
)
```

```
/usr/local/lib/python3.7/dist-packages/torchtext/data/iterator.py:48: UserWarning: Bucket
warnings.warn('{}' class will be retired soon and moved to torchtext.legacy. Please
```

Вы можете использовать Conv2d с `in_channels=1, kernel_size=(kernel_sizes[0], emb_dim))` или Conv1d с `in_channels=emb_dim, kernel_size=kernel_size[0]`. Но хорошенько подумайте над shape в обоих случаях.

```
class CNN(nn.Module):
    def __init__(
        self,
        vocab_size,
        emb_dim,
        out_channels,
        kernel_sizes,
        dropout=0.5,
    ):
        super().__init__()

        self.embedding = nn.Embedding(vocab_size, emb_dim)
        self.conv_0 = nn.Conv1d(emb_dim, out_channels, kernel_size=kernel_sizes[0], padding=1)
        self.conv_1 = nn.Conv1d(emb_dim, out_channels, kernel_size=kernel_sizes[1], padding=1)
        self.conv_2 = nn.Conv1d(emb_dim, out_channels, kernel_size=kernel_sizes[2], padding=1)
        self.fc = nn.Linear(len(kernel_sizes) * out_channels, 1)

        self.dropout = nn.Dropout(dropout)
```

```

def forward(self, text):

    embedded = self.embedding(text) # [batch_size, seq_len, embed_dim]
    embedded = embedded.permute(0, 2, 1) # may be reshape here

    convded_0 = F.relu(self.conv_0(embedded)) # may be reshape here
    convded_1 = F.relu(self.conv_1(embedded)) # may be reshape here
    convded_2 = F.relu(self.conv_2(embedded)) # may be reshape here

    pooled_0 = F.max_pool1d(convded_0, convded_0.shape[2]).squeeze(2)
    pooled_1 = F.max_pool1d(convded_1, convded_1.shape[2]).squeeze(2)
    pooled_2 = F.max_pool1d(convded_2, convded_2.shape[2]).squeeze(2)

    cat = self.dropout(torch.cat((pooled_0, pooled_1, pooled_2), dim=1))

    return self.fc(cat)

kernel_sizes = [3, 4, 5]
vocab_size = len(TEXT.vocab)
out_channels=64
dropout = 0.5
dim = 300
patience=2

model = CNN(vocab_size=vocab_size, emb_dim=dim, out_channels=out_channels,
            kernel_sizes=kernel_sizes, dropout=dropout)

model.to(device)

CNN(
  (embedding): Embedding(202268, 300)
  (conv_0): Conv1d(300, 64, kernel_size=(3,), stride=(2,), padding=(1,))
  (conv_1): Conv1d(300, 64, kernel_size=(4,), stride=(2,), padding=(1,))
  (conv_2): Conv1d(300, 64, kernel_size=(5,), stride=(2,), padding=(1,))
  (fc): Linear(in_features=192, out_features=1, bias=True)
  (dropout): Dropout(p=0.5, inplace=False)
)

opt = torch.optim.Adam(model.parameters())
loss_func = nn.BCEWithLogitsLoss()

max_epochs = 30

Обучите!

import numpy as np

min_loss = np.inf

cur_patience = 0

```

```

for epoch in range(1, max_epochs + 1):

    CNN_all_preds_train = []
    CNN_all_label_train = []
    CNN_all_preds_val = []
    CNN_all_label_val = []
    CNN_correct_train = 0
    CNN_correct_val = 0
    CNN_num_obj_train = 0
    CNN_num_obj_val = 0
    train_loss = 0.0

    model.train()
    pbar = tqdm(enumerate(train_iter), total=len(train_iter), leave=False)
    pbar.set_description(f"Epoch {epoch}")
    for it, batch in pbar:
        opt.zero_grad()
        emb = batch.text.to(device)
        # print("emb Shape: ", emb.shape)
        # print("emb: ", emb)
        label = torch.unsqueeze(batch.label, 1).to(device)
        # print("label Shape: ", label.shape)
        # print("label: ", label)
        preds = model(emb)
        loss = loss_func(preds, label)
        loss.backward()
        train_loss += loss
        opt.step()

        preds_r = np.where(preds.cpu() > 0, 1, 0)
        label_r = np.where(label.cpu() > 0, 1, 0)
        c = preds_r == label_r
        CNN_correct_train += np.count_nonzero(c)
        CNN_num_obj_train += len(batch.label)

        trn = np.sum(preds_r, axis=1)
        lbl = np.sum(label_r, axis=1)
        trn = trn.tolist()
        lbl = lbl.tolist()
        CNN_all_preds_train.append(trn)
        CNN_all_label_train.append(lbl)
        #YOUR CODE GOES HERE

    train_loss /= len(train_iter)
    val_loss = 0.0
    model.eval()
    pbar = tqdm(enumerate(val_iter), total=len(val_iter), leave=False)
    pbar.set_description(f"Epoch {epoch}")
    for it, batch in pbar:
        emb = batch.text.to(device)
        #len_emb = batch.text[1].cpu()
        label = torch.unsqueeze(batch.label, 1).to(device)
        preds = model(emb)

```

```

predsr = np.where(preds.cpu() > 0, 1, 0)
labelr = np.where(label.cpu() > 0, 1, 0)
c = predsr == labelr
CNN_correct_val += np.count_nonzero(c)
CNN_num_obj_val += len(batch.label)

val_loss += loss_func(preds.data, label)

val = np.sum(predsr, axis=1)
lbl_val = np.sum(labelr, axis=1)
val = val.tolist()
lbl_val = lbl_val.tolist()
CNN_all_preds_val.append(val)
CNN_all_label_val.append(lbl_val)
# YOUR CODE GOES HERE

val_loss /= len(val_iter)
if val_loss < min_loss:
    min_loss = val_loss
    best_model = model.state_dict()
else:
    cur_patience += 1
    if cur_patience == patience:
        cur_patience = 0
        break

print('Epoch: {}, Training Loss: {}, Validation Loss: {}, Training accuracy: {}, Valid
model.load_state_dict(best_model)

/usr/local/lib/python3.7/dist-packages/torchtext/data/batch.py:23: UserWarning: Batch
warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please
Epoch: 1, Training Loss: 0.6645534634590149, Validation Loss: 0.5257830023765564, Tra
Epoch: 2, Training Loss: 0.5165377259254456, Validation Loss: 0.47080883383750916, Tr
Epoch: 3, Training Loss: 0.4343946874141693, Validation Loss: 0.4132627248764038, Tra
Epoch: 4, Training Loss: 0.3629172146320343, Validation Loss: 0.384289413690567, Trai
Epoch: 5, Training Loss: 0.2830201983451843, Validation Loss: 0.36714044213294983, Tr
Epoch: 6, Training Loss: 0.21790607273578644, Validation Loss: 0.3537408113479614, Tr
Epoch: 7, Training Loss: 0.14971472322940826, Validation Loss: 0.35701942443847656, 1
<All keys matched successfully>

```

Посчитайте f1-score вашего классификатора.

Ответ:

```

CNN_flat_preds_list = [item for sublist in CNN_all_preds_val for item in sublist]
CNN_flat_label_list = [item for sublist in CNN_all_label_val for item in sublist]

```

```
f1_score(CNN_flat_label_list, CNN_flat_preds_list, average='weighted')
```

```
0.8467762355182719
```

▼ Интерпретируемость

Посмотрим куда смотрит наша модель. Достаточно запустить код ниже

```
!pip install -q captum
```

```
|████████████████████████████████████████████████████████████████████████████████| 4.4MB 3.8MB/s
```

```
from captum.attr import LayerIntegratedGradients, TokenReferenceBase, visualization
```

```
PAD_IND = TEXT.vocab.stoi['pad']
```

```
token_reference = TokenReferenceBase(reference_token_idx=PAD_IND)
```

```
lig = LayerIntegratedGradients(model, model.embedding)
```

```
def forward_with_softmax(inp):
    logits = model(inp)
    return torch.softmax(logits, 0)[0][1]
```

```
def forward_with_sigmoid(input):
    return torch.sigmoid(model(input))
```

```
# accumulate couple samples in this array for visualization purposes
```

```
vis_data_records_ig = []
```

```
def interpret_sentence(model, sentence, min_len = 7, label = 0):
```

```
    model.eval()
    text = [tok for tok in TEXT.tokenize(sentence)]
    if len(text) < min_len:
        text += ['pad'] * (min_len - len(text))
    indexed = [TEXT.vocab.stoi[t] for t in text]
```

```
    model.zero_grad()
```

```
    input_indices = torch.tensor(indexed, device=device)
    input_indices = input_indices.unsqueeze(0)
```

```
    # input_indices dim: [sequence_length]
    seq_length = min_len
```

```
    # predict
    pred = forward_with_sigmoid(input_indices).item()
    pred_ind = round(pred)
```

```
    # generate reference indices for each sample
    reference_indices = token_reference.generate_reference(seq_length, device=device).unsq
```

```
    # compute attributions and approximation delta using layer integrated gradients
    attributions_ig, delta = lig.attribute(input_indices, reference_indices, \
                                           n_steps=5000, return_convergence_delta=True)
```

```
    print('pred: ', LABEL.vocab.itos[pred_ind], '(', '%.2f'%pred, ')', ', delta: ', abs(de
```

```

add_attributions_to_visualizer(attributions_ig, text, pred, pred_ind, label, delta, vi

def add_attributions_to_visualizer(attributions, text, pred, pred_ind, label, delta, vis_d
    attributions = attributions.sum(dim=2).squeeze(0)
    attributions = attributions / torch.norm(attributions)
    attributions = attributions.cpu().detach().numpy()

# storing couple samples in an array for visualization purposes
vis_data_records.append(visualization.VisualizationDataRecord(
    attributions,
    pred,
    LABEL.vocab.itos[pred_ind],
    LABEL.vocab.itos[label],
    LABEL.vocab.itos[1],
    attributions.sum(),
    text,
    delta))

interpret_sentence(model, 'It was a fantastic performance !', label=1)
interpret_sentence(model, 'Best film ever', label=1)
interpret_sentence(model, 'Such a great show!', label=1)
interpret_sentence(model, 'It was a horrible movie', label=0)
interpret_sentence(model, 'I\'ve never watched something as bad', label=0)
interpret_sentence(model, 'It is a disgusting movie!', label=0)

    pred: pos ( 1.00 ) , delta: tensor([0.0003], device='cuda:0', dtype=torch.float64)
    pred: pos ( 0.97 ) , delta: tensor([8.6919e-05], device='cuda:0', dtype=torch.float
    pred: pos ( 1.00 ) , delta: tensor([6.2853e-06], device='cuda:0', dtype=torch.float
    pred: neg ( 0.01 ) , delta: tensor([0.0001], device='cuda:0', dtype=torch.float64)
    pred: neg ( 0.02 ) , delta: tensor([9.2424e-05], device='cuda:0', dtype=torch.float
    pred: pos ( 0.94 ) , delta: tensor([8.4931e-06], device='cuda:0', dtype=torch.float

```

Попробуйте добавить свои примеры!

```

print('Visualize attributions based on Integrated Gradients')
visualization.visualize_text(vis_data_records_ig)

```

Visualize attributions based on Integrated Gradients

Legend: ☐ Negative ☐ Neutral ☐ Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (1.00)	pos	0.60	It was a fantastic performance ! pad
pos	pos (0.97)	pos	0.14	Best film ever pad pad pad pad
pos	pos (1.00)	pos	0.67	Such a great show! pad pad pad
neg	neg (0.01)	pos	-1.00	It was a horrible movie pad pad

▼ Эмбэдинги слов

Вы ведь не забыли, как мы можем применить знания о word2vec и GloVe. Давайте попробуем!

```

True      Predicted      Attribution      Attribution
TEXT.build_vocab(trn, vectors=GloVe(name='6B', dim=300))# YOUR CODE GOES HERE
# подсказка: один из импортов пока не использовался, быть может он нужен в строке выше :)
LABEL.build_vocab(trn)

word_embeddings = TEXT.vocab.vectors

kernel_sizes = [3, 4, 5]
vocab_size = len(TEXT.vocab)
dropout = 0.5
dim = 300

train, tst = datasets.IMDB.splits(TEXT, LABEL)
trn, vld = train.split(random_state=random.seed(SEED))

device = "cuda" if torch.cuda.is_available() else "cpu"

train_iter, val_iter, test_iter = BucketIterator.splits(
    (trn, vld, tst),
    batch_sizes=(128, 256, 256),
    sort=False,
    sort_key= lambda x: len(x.src),
    sort_within_batch=False,
    device=device,
    repeat=False,
)

model = CNN(vocab_size=vocab_size, emb_dim=dim, out_channels=64,
            kernel_sizes=kernel_sizes, dropout=dropout)

word_embeddings = TEXT.vocab.vectors

prev_shape = model.embedding.weight.shape

model.embedding.weight = torch.nn.Parameter(word_embeddings) # инициализируйте эмбэдинги
assert prev_shape == model.embedding.weight.shape

```

```
assert prev_shape == model.embedding.weight.shape
model.to(device)
```

```
opt = torch.optim.Adam(model.parameters())
```

Вы знаете, что делать.

```
import numpy as np
```

```
patience=3
```

```
min_loss = np.inf
```

```
cur_patience = 0
```

```
for epoch in range(1, max_epochs + 1):
```

```
    EMBCNN_all_preds_train = []
    EMBCNN_all_label_train = []
    EMBCNN_all_preds_val = []
    EMBCNN_all_label_val = []
    EMBCNN_correct_train = 0
    EMBCNN_correct_val = 0
    EMBCNN_num_obj_train = 0
    EMBCNN_num_obj_val = 0
    train_loss = 0.0
```

```
    model.train()
```

```
    pbar = tqdm(enumerate(train_iter), total=len(train_iter), leave=False)
```

```
    pbar.set_description(f"Epoch {epoch}")
```

```
    for it, batch in pbar:
```

```
        opt.zero_grad()
        emb = batch.text.to(device)
        label = torch.unsqueeze(batch.label, 1).to(device)
        preds = model(emb)
        loss = loss_func(preds, label)
        loss.backward()
        train_loss += loss
        opt.step()
```

```
        preds_r = np.where(preds.cpu() > 0, 1, 0)
        label_r = np.where(label.cpu() > 0, 1, 0)
        c = preds_r == label_r
        EMBCNN_correct_train += np.count_nonzero(c)
        EMBCNN_num_obj_train += len(batch.label)
```

```
        trn = np.sum(preds_r, axis=1)
        lbl = np.sum(label_r, axis=1)
        trn = trn.tolist()
        lbl = lbl.tolist()
        EMBCNN_all_preds_train.append(trn)
        EMBCNN_all_label_train.append(lbl)
        #YOUR CODE GOES HERE
```



```

train_loss /= len(train_iter)
val_loss = 0.0
model.eval()
pbar = tqdm(enumerate(val_iter), total=len(val_iter), leave=False)
pbar.set_description(f"Epoch {epoch}")
for it, batch in pbar:

    emb = batch.text.to(device)
    label = torch.unsqueeze(batch.label, 1).to(device)
    preds = model(emb)

    preds_r = np.where(preds.cpu() > 0, 1, 0)
    label_r = np.where(label.cpu() > 0, 1, 0)
    c = preds_r == label_r
    EMBCNN_correct_val += np.count_nonzero(c)
    EMBCNN_num_obj_val += len(batch.label)

    val_loss += loss_func(preds.data, label)

    val = np.sum(preds_r, axis=1)
    lbl_val = np.sum(label_r, axis=1)
    val = val.tolist()
    lbl_val = lbl_val.tolist()
    EMBCNN_all_preds_val.append(val)
    EMBCNN_all_label_val.append(lbl_val)
    # YOUR CODE GOES HERE
val_loss /= len(val_iter)
if val_loss < min_loss:
    min_loss = val_loss
    best_model = model.state_dict()
else:
    cur_patience += 1
    if cur_patience == patience:
        cur_patience = 0
        break

print('Epoch: {}, Training Loss: {}, Validation Loss: {}, Training accuracy: {}, Valid
model.load_state_dict(best_model)

```

Посчитайте f1-score вашего классификатора.

Ответ:

```

EMBCNN_flat_preds_list = [item for sublist in EMBCNN_all_preds_val for item in sublist]
EMBCNN_flat_label_list = [item for sublist in EMBCNN_all_label_val for item in sublist]

```

```
f1_score(EMBCNN_flat_label_list, EMBCNN_flat_preds_list, average='weighted')
```

```
0.8706160394010135
```

Проверим насколько все хорошо!

```
PAD_IND = TEXT.vocab.stoi['pad']
```

```
token_reference = TokenReferenceBase(reference_token_idx=PAD_IND)
```

```
lig = LayerIntegratedGradients(model, model.embedding)
```

```
vis_data_records_ig = []
```

```
interpret_sentence(model, 'It was a fantastic performance !', label=1)
```

```
interpret_sentence(model, 'Best film ever', label=1)
```

```
interpret_sentence(model, 'Such a great show!', label=1)
```

```
interpret_sentence(model, 'It was a horrible movie', label=0)
```

```
interpret_sentence(model, 'I\'ve never watched something as bad', label=0)
```

```
interpret_sentence(model, 'It is a disgusting movie!', label=0)
```

```
pred: pos ( 0.91 ) , delta: tensor([0.0001], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.00 ) , delta: tensor([6.2463e-06], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.03 ) , delta: tensor([0.0002], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.00 ) , delta: tensor([6.8991e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.09 ) , delta: tensor([0.0001], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.00 ) , delta: tensor([7.4932e-05], device='cuda:0', dtype=torch.float64)
```

```
print('Visualize attributions based on Integrated Gradients')
```

```
visualization.visualize_text(vis_data_records_ig)
```

Visualize attributions based on Integrated Gradients

Legend: ☐ Negative ☐ Neutral ☐ Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (0.91)	pos	1.89	It was a fantastic performance ! pad
pos	neg (0.00)	pos	1.57	Best film ever pad pad pad pad
pos	neg (0.03)	pos	1.69	Such a great show! pad pad pad
neg	neg (0.00)	pos	0.01	It was a horrible movie pad pad
neg	neg (0.09)	pos	1.29	I've never watched something as bad pad
neg	neg (0.00)	pos	0.09	It is a disgusting movie! pad pad

Legend: ☐ Negative ☐ Neutral ☐ Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (0.91)	pos	1.89	It was a fantastic performance ! pad
pos	neg (0.00)	pos	1.57	Best film ever pad pad pad pad
pos	neg (0.03)	pos	1.69	Such a great show! pad pad pad
neg	neg (0.00)	pos	0.01	It was a horrible movie pad pad
				I've never watched something as bad

